

P33 Building an Internet Router - Documentation Part 1

C. T. Bashford-Chuchla (ctb33)

Contents

1	Code Flow	1
1.1	ARP Packets	1
1.2	IP Packets	2
1.2.1	ICMP Packets	3
1.2.2	PWOSPF Packets	3
1.3	Router	5
2	Design Decisions	6
2.1	Packets	6
2.2	ARP Protocol	6
2.3	Routing Protocol	6
2.3.1	Storage	6

1 Code Flow

Packets enter my code in the `router_handle_packet (packet_info_t *pi)` function. Of which packets separate into two directions, one for ARP packets and the other for IP packets.

1.1 ARP Packets

All incoming ARP packets are sent through to `handle_ARP_packet(packet_info_t *pi)`. All functions that deal with ARP packets are contained in the *arp.c* file.

1. `handle_ARP_packet(packet_info_t *pi)`
This function deals with the processing of the incoming ARP packets. It is split into two, first part for requests and the second for replies. The request part first stores the mac address of the sender in the arp table, and replies by modifying the packet (changes the OP code to 2, swapping the sender and target details, and sets the sender mac address). The reply part saves the sender mac address in the ARP table and then checks if there are any pending packets that need the ARP information just received (see TODO: ARP packets in Design Decisions) and sends them if necessary.
2. `handle_not_responding_to_arp (byte *payload, unsigned len)`
This function generates a not responding to arp packet with the given packet. Most of this function is achieved by calling `generate_response_ICMP_packet(packet_info_t *pi, int type, int code)` where type is 3 and code is 1.
3. `generate_pending_ARP_thread()`
This function is constantly being run in the ARP thread (see TODO: ARP packets in Design Decisions). It has two main functions, first of which is to send ARP packets that are in the `pending_arp` array

and secondly to send not responding to arp packets(see above) when there is no response after 5 arp requests sent.

4. `send_ARP_request(addr_ip_t ip, int num)`
This function sends an ARP request to the specified ip address parameter. It is used in the ARP thread by the `generate_pending_ARP_thread()` function.
5. `router_add_arp_entry (router_t *router, addr_mac_t mac, addr_ip_t ip, bool dynamic)`
This function adds the arp entry given into the arp table. It also sends the same ARP entry to the hardware, if running in hardware mode.
6. `router_delete_arp_entry (router_t *router, addr_ip_t ip)`
This function deletes an arp entry specified by the ip parameter. After deleting the entry, it also shifts all entries after it up by one, so there are no gaps in the table. This is all reflected in the hardware ARP table, by sending the same ARP entries to the hardware.
7. `router_delete_all_arp_entries (router_t *router, bool dynamic)`
This function deletes all ARP entries that are specified by the dynamic boolean parameter. This is achieved by deleting the entries and shifting succeeding entries in one loop through the arp entries. The hardware ARP table is updated by sending all the ARP entries to the hardware at the end.
8. `router_find_arp_entry (router_t *router, addr_ip_t ip)`
This function finds the ARP entry in the ARP table specified by the ip parameter. The function returns a pointer to the ip, mac pair if found, otherwise returns NULL.

1.2 IP Packets

All incoming IP packets are sent through to `handle_IPv4_packet(packet_info_t *pi)`. All ICMP packets are sent through to `handle_ICMP_packet(packet_info_t *pi)` (see ICMP section TODO), TCP packets to `handle_TCP_packet(packet_info_t *pi)` (contained in *ip.c*) and PWOSPF packets to `handle_PWOSPF_packet(packet_info_t *pi)`. All functions that deal with IP packets are contained in the *ip.c* file.

1. `handle_IPv4_packet(packet_info_t *pi)`
This function is called with all incoming IP packets. It first of all checks that there are no options in the packet and then checks the checksum. If the packet has a destination of one of the interfaces or the OSPF IP address then it is dealt by the router directly using the functions specified above, otherwise the TTL is decreased by one (and recompute checksum) and sends the packet towards its destination.
2. `handle_TCP_packet(packet_info_t *pi)`
All TCP packets that are intended for one of the router's interfaces are called with this function. The ethernet header is stripped off and the rest of the packet is sent to the `sr_transport_input (uint8_t * packet)` function.
3. `calc_checksum(byte *header, int len)`
This function just computes the checksum of the header specified with the length given.
4. `add_IPv4_header(uint8_t * payload, uint8_t proto, uint32_t src, uint32_t dest, int len)`
This function takes a packet and puts a IPv4 header on it with the specified protocol, source and destination IP. A new packet is malloc'ed with 20 bytes longer length and the function returns the new pointer to the packet.

5. `handle_no_route_to_host (packet_info_t *pi)`

This function is called when there is no route to the host in the routing table. A ICMP packet is sent to the source specifying this. Most of this function is achieved by calling `generate_response_ICMP_packet(packet_info_t *pi, int type, int code)` where type is 3 and code is 0.

1.2.1 ICMP Packets

All incoming ICMP packets are sent through to `handle_ICMP_packet(packet_info_t *pi)`. All functions that deal with ICMP packets are contained in the *my_icmp.c* file.

1. `handle_ICMP_packet(packet_info_t *pi)`

This function is called with all incoming ICMP packets. The checksum for the ICMP packet is calculated first of all. Then if the packet is an echo request, the type is changed to 0 and the checksum is recalculated. Otherwise, if the packet is a echo reply then `handle_ping_reply (packet_info_t *pi)` is called (see below). All echo request packets (after type is set to 0) are sent in the tail part of the `handle_IPv4_packet(packet_info_t *pi)` function after this packet returns.

2. `generate_response_ICMP_packet(packet_info_t *pi, int type, int code)`

This function creates a ICMP packet with IPv4 header with the specified type and code. The function creates a new packet and sets it in the `packet_info_t` structure. The destination is computed by the source of the inputted packet.

3. `void send_ping(router_t *router, addr_ip_t dest_ip, addr_ip_t src_ip, uint16_t id, uint16_t count)`

This function sends a ping to the specified destination IP address. This function is called only by `cli_ping_request (router_t *rtr, int fd, addr_ip_t ip)` in *cli_ping.c*. A IPv4 is added to the ping packet in this function and is sent to the correct interface.

4. `handle_ping_reply (packet_info_t *pi)`

This function deals with an incoming ping reply packet. This is called by `handle_ICMP_packet(packet_info_t *pi)`. The packet is sent to the cli via the function `cli_ping_handle_reply (addr_ip_t ip, uint16_t seq)`.

1.2.2 PWOSPF Packets

All incoming PWOSPF packets are sent through to `handle_PWOSPF_packet(packet_info_t *pi)`. All functions that deal with PWOSPF packets are contained in the *routing.c* file. This includes computing the routing table.

1. `handle_PWOSPF_packet(packet_info_t *pi)`

This function is called with all incoming PWOSPF packets. First of all, there is a check to see when the router is using the OSPF protocol currently. If it is not then all packets are dropped. If the OSPF protocol is running then the version value is checked that it is 2, the authentication is checked that it is 0, the checksum is checked and finally that the area id in the packet is the same as the router's area id. If any of these checks fail then the packet is dropped. The packet goes into two paths, one for HELLO packets and the other for Link State Update packets. For the HELLO packets, the subnet_mask and helloint are checked against the interface the packet came in on. If there is a mismatch the packet is dropped. The neighbours are added to the neighbor linked list, if they are not already on it. Otherwise, the last time recieved counter in the corresponding neighbor entry is updated and so is the link state update for the neighbor in the link state database for the current router. Lastly, if a new neighbor was added to the linked list, then a link state is added to the database for the current router. Now, if the packet is a Link State Update, the router id is checked that is for this router. If there is no database

entry for this router, then one is created with the specified link state adverts and the routing table is recomputed. Otherwise, the sequence number is checked that it is different to the previous and that the contents is different to the last link state update received from this router. If either check fails then the packet is dropped. First of all the link state adverts are compared with the current link state adverts in the database for the source router. If there is no change then the last time received counter are updated and it finishes. If there is a change in the link states adverts in the database or the link state packet has come from a router not in the database then the new advertisements link states are added (with replacement) and the routing table is updated by calling `update_routing_table ()`. Finally, if the packet hasn't been dropped, the packet's TTL is decreased by one (with recomputing the checksum) and forwarded to all neighbors.

2. `print_database ()`
This function prints out the database of stored link state updates for all routes.
3. `update_routing_table ()`
This function first finds all subnets in the link state database and calculates a distance to each subnet. For all reachable subnets a routing entry is added. This is explained in more detail in the Design Decisions TODO section.
4. `send_HELLO_packet(interface_t * intf)`
This function is used to send HELLO packets to the specified interface. It computes the desired packet by the interface information and sends it if there is a neighbor connected.
5. `send_LSU_packet(unsigned seq_no)`
This function is used to send link state update packets with the specified sequence number to interfaces with routers as neighbors. The function computes the link states adverts from the neighbors on each interface and puts them into a packet to be forwarded to each interface.
6. `generate_HELLO_thread()`
This function is constantly being run in the HELLO thread (see TODO: HELLO packets in Design Decisions). It has three main functions, first of which is to send HELLO packets to each connected neighbor every `helloint` seconds, secondly to check for neighbor timeouts (and if there is to update the routing table and send a new link state update packet to all interfaces) and lastly to send link state update packets to all interfaces every `lsuint` seconds or if there a link state change has been detected.
7. `router_add_route (router_t * router , addr_ip_t prefix , addr_ip_t next_hop , addr_ip_t subnet_mask , const char *intf_name , bool dynamic)`
This function adds the specified route to the routing table in the correct place (sorted by dynamic and then network mask), it does this by shifting the succeeding routes down by one. The same changes are also reflected in the hardware routing table.
8. `router_find_route_entry (router_t *router , addr_ip_t dest , addr_ip_t gw , addr_ip_t mask , const char *intf_name)`
This function finds the entry in the routing table specified by the parameters. The function returns a pointer to the routing entry if found, otherwise returns NULL.
9. `router_delete_route_entry (router_t *router , addr_ip_t dest , addr_ip_t gw , addr_ip_t mask , const char *intf_name)`
This function deletes the route in the routing table specified by the parameters. The function returns whether it is successful in finding a matching entry. All changes are reflected in the hardware routing table.

10. `router_delete_all_route_entries (router_t *router , bool dynamic)`
This function deletes all routes in the routing table that matches the dynamic boolean parameter. All changes are reflected in the hardware routing table.
11. `sr_read_routes_from_file (router_t * router , const char* filename)`
This function reads in the routing file and adds all routes as static in the routing table.
12. `database_find_link (database_entry_t *database_entry , uint32_t router_id , uint32_t subnet_no)`
This function finds the link in the database entry specified by the parameters. The function returns a pointer to the link if found, otherwise returns NULL.
13. `router_add_link_to_database_entry (router_t *router , database_entry_t *database_entry , link_t *link_to_add)`
This function adds the link to the database entry specified by the database_entry and link_to_add parameters respectively. If the link already exists in the database, then the last time is updated.
14. `router_remove_link_from_database_entry (router_t *router , database_entry_t *database_entry , uint32_t router_id)`
This function removes the link in the database entry specified by the database_entry and router_id parameters respectively. The function returns whether the link was found in the database.
15. `router_add_database_entry (router_t * router , uint32_t router_id , link_t link [], unsigned num_links, uint16_t seq_no , byte *packet, unsigned len)`
This function adds all the links in the specified link array to the database entry specified by the router id. If the links exceed 10 then an error is thrown. Each link is copied and thus the array parameter will not need to be kept in memory.
16. `router_find_database_entry_position (router_t *router , uint32_t router_id)`
This function find the position of the database entry specified by the router id. This will return the integer in which the router id entry is in the database array.

1.3 Router

This section reflects the changes in *sr_router.c*. The main purpose of the functions added was to send packets and to setup the hardware interface table.

1. `send_packet(byte *payload, uint32_t src, uint32_t dest, int len, bool is_arp_packet , bool is_hello_packet)`
This function sends a packet specified by payload to the destination specified. It first checks the destination is one of the interfaces or the OSPF IP address. If it is not, a no route to host packet is generated. Once passing this test, the packet is sent to the found interface by `send_packet_intf (...)` (see below).
2. `send_packet_intf (interface_t *intf , byte *payload, uint32_t src, uint32_t dest, int len, bool is_arp_packet , bool is_hello_packet)`
This function sends the specified packet on the interface specified by the argument intf. The ethernet packet is added to the packet, in doing so it checks for an ARP entry, and sends the packet. If there is no ARP entry for the next hop, then the packet is queued and an ARP request is sent.
3. `setup_interface_registers (router_t * router , int intf_num)`
This function sends the interface mac address of the specified interface to the hardware interface table.

2 Design Decisions

2.1 Packets

In most cases when responding to a packet, the code will modify the incoming packet in order to responding to the packet. An exception to this rule is when generating response ICMP packets, as this will often need to be longer than the incoming packet. The decision was chosen because meant less memory need to be created and that less packet fields would need to be generated as they can be reused. This also means, that a lot of the code has a tail. In other words, in the `handle_IPv4_packet()` function the packet is sent in the tail of the function after the `handle_ICMP_packet()` finishes.

2.2 ARP Protocol

In order to deal with pending ARP requests, a seperate thread was used. The thread constantly runs the `generate_pending_ARP_thread()` function. The approach uses the thread to check for packets in the pending ARP queue and send up to five ARP requests if necessary, and more over generate a not responding to ARP request response. This was chosen over starting a seperate thread for each pending ARP request, as running too many threads caused some threads to not get enough CPU time to send the packets in a timely fashion. In order to keep the threads safe, locks were used when ever modifying the pending arp table.

2.3 Routing Protocol

2.3.1 Storage

The routing protcol required two elements to be recorded, the neighbors in which a HELLO packet was recieved and the link state advertisments recieved. For the neighbors, a linked list was created on each interface was all neighbors that send to HELLO.